

Poster [G3]



Deep Spatio-Temporal Random Fields for Efficient Video Segmentation

Siddhartha Chandra, Camille Couprie & Iasonas Kokkinos



Video G-CRFs: Motivation



Frame-by-Frame Prediction

Video G-CRFs: Motivation



Frame-by-Frame Prediction



Video G-CRF

Video G-CRFs: Applications



Baseline

Video G-CRF

Video Instance Segmentation



Baseline

Video G-CRF

Video Semantic Segmentation




Baseline


Video G-CRF

Instance Tracking


Deep Spatio-Temporal G-CRF




Decoupled Predictions




Deep Spatio-Temporal G-CRF




Sparse Graphical Model





Deep Spatio-Temporal G-CRF




Dense Graphical Model




Deep Spatio-Temporal G-CRF




Dense Graphical Model
(No coupling of predictions)





Deep Spatio-Temporal G-CRF




Dense Spatio-Temporal
Graphical Model



Deep Spatio-Temporal G-CRF



Dense Spatio-Temporal
Graphical Model




Deep Spatio-Temporal G-CRF

Gaussian-CRFs:

$$E(\mathbf{x}) = \underbrace{\frac{1}{2} \mathbf{x}^T A \mathbf{x}}_{\text{pairwise}} - \underbrace{B^T \mathbf{x}}_{\text{unary}}$$

$$T_{u,v,p_i,p_j}(l_m, l_n) = \langle \mathcal{T}_{u,p_i}^{l_m}, \mathcal{T}_{v,p_j}^{l_n} \rangle$$



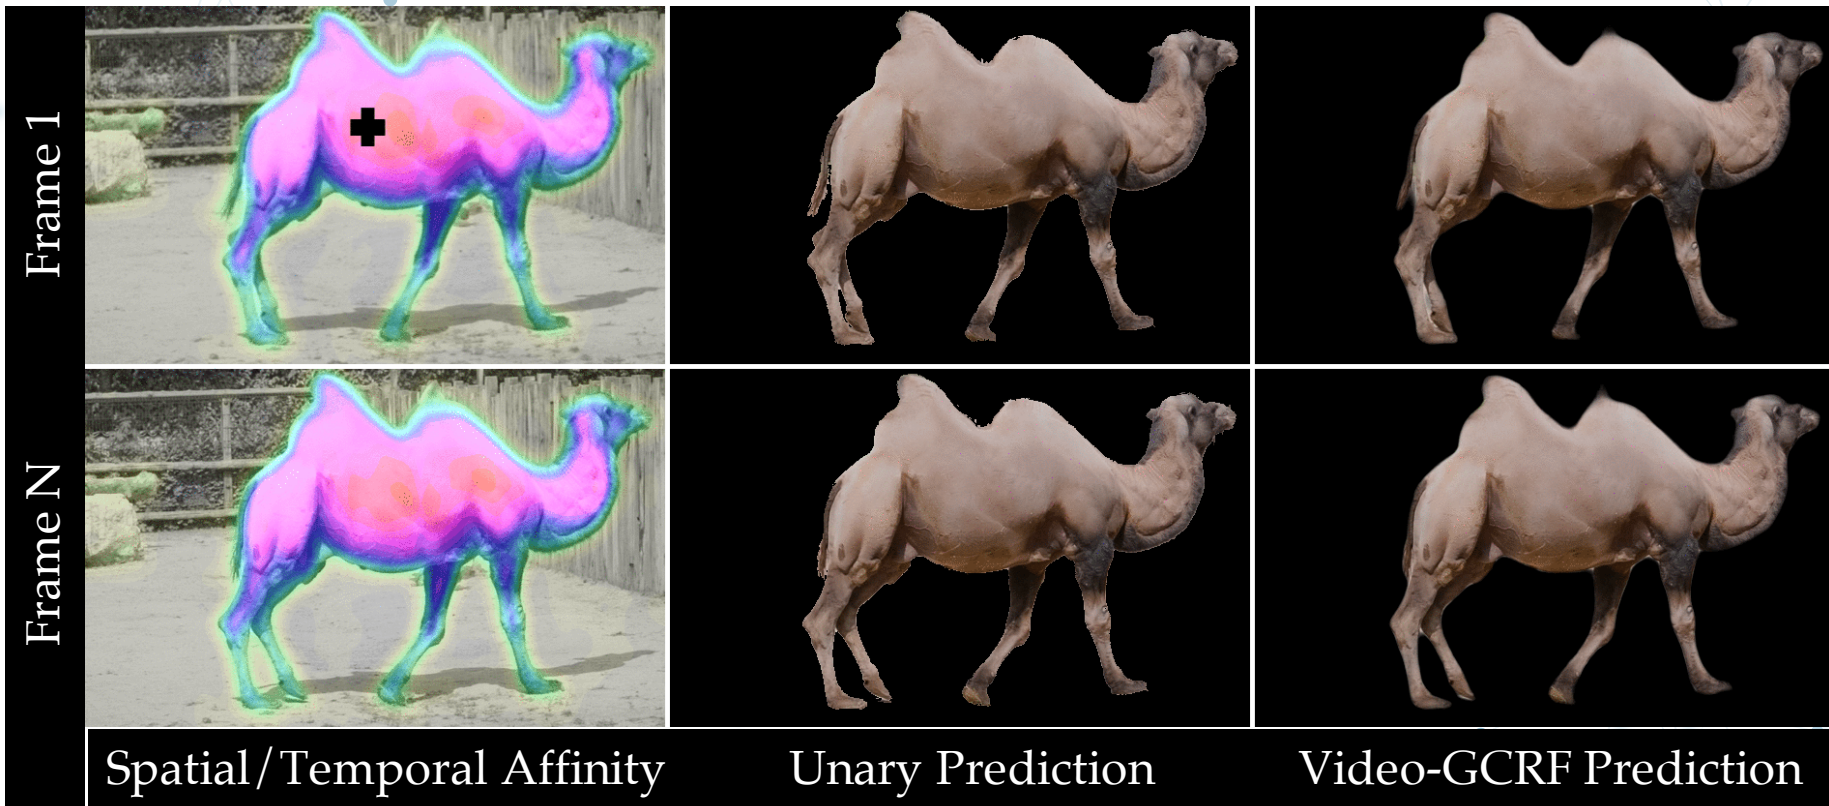
$$E(\mathbf{x}) = - \underbrace{\sum_i b_i \mathbf{x}_i}_{\text{unary}} + \frac{1}{2} \underbrace{\sum_{i,j} \mathbf{x}_i \mathbf{x}_j a_{i,j}}_{\text{pairwise}}$$

$$A_{v,p_i,p_j}(l_m, l_n) = \langle \mathcal{A}_{v,p_i}^{l_m}, \mathcal{A}_{v,p_j}^{l_n} \rangle$$


Inference:

$$A \mathbf{x} = B$$

Pairwise Affinities: Instance Tracking

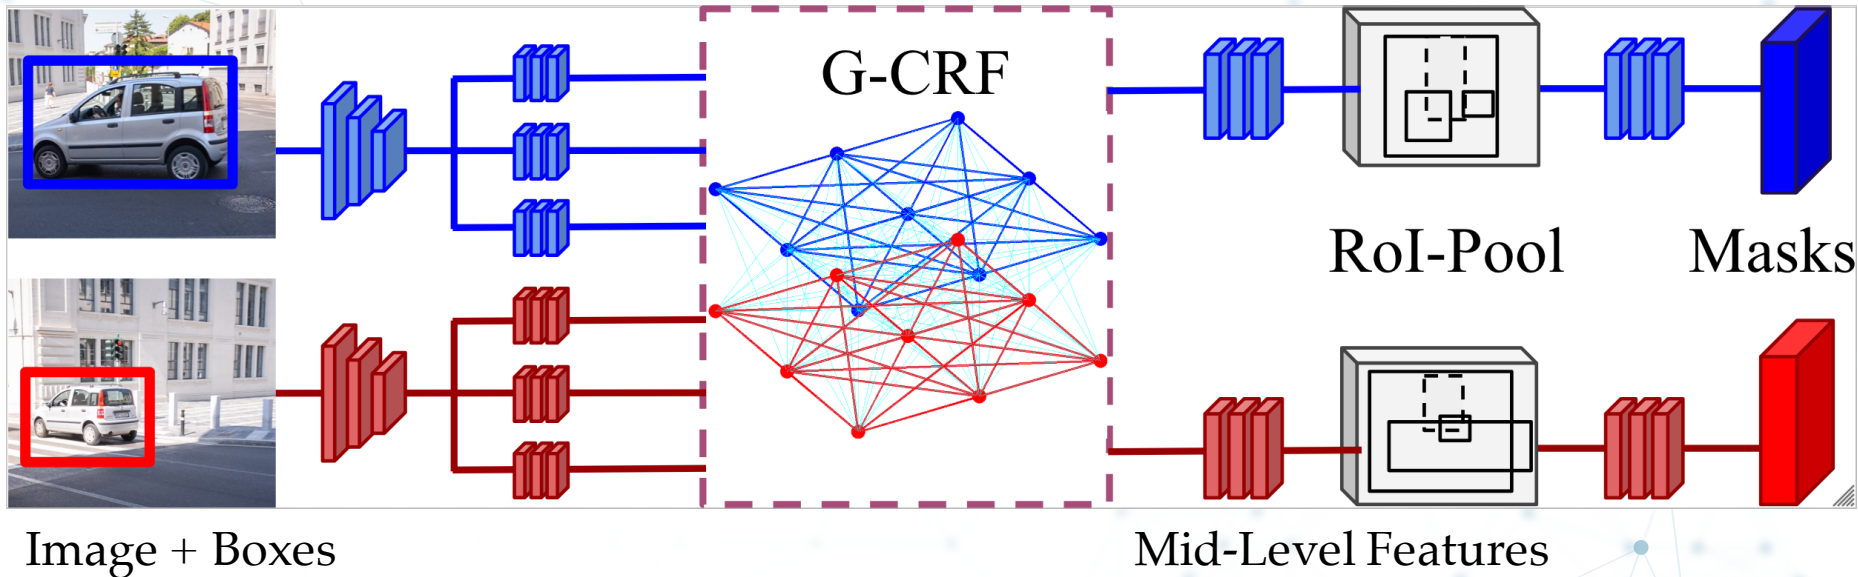


Video G-CRFs: Overview




Video G-CRFs: Instance Segmentation

Extension of Mask RCNN [He 2017]



Video G-CRFs: Instance Tracking



Frame-by-Frame Prediction



Video G-CRF

Video G-CRFs: Semantic Segmentation

